

Company_Bankruptcy

July 17, 2023

1 Company Bankruptcy Prediction

Bankruptcy data from the Taiwan Economic Journal for the years 1999–2009

1.1 Goal

how many firms are bankrupt, and how many are not.

2 Attribute Information

Updated column names and description to make the data easier to understand (Y = Output feature, X = Input features)

Y - Bankrupt?: Class label

X1 - ROA(C) before interest and depreciation before interest: Return On Total Assets(C)

X2 - ROA(A) before interest and % after tax: Return On Total Assets(A)

X3 - ROA(B) before interest and depreciation after tax: Return On Total Assets(B)

X4 - Operating Gross Margin: Gross Profit/Net Sales

X5 - Realized Sales Gross Margin: Realized Gross Profit/Net Sales

X6 - Operating Profit Rate: Operating Income/Net Sales

X7 - Pre-tax net Interest Rate: Pre-Tax Income/Net Sales

X8 - After-tax net Interest Rate: Net Income/Net Sales

X9 - Non-industry income and expenditure/revenue: Net Non-operating Income Ratio

X10 - Continuous interest rate (after tax): Net Income-Exclude Disposal Gain or Loss/Net Sales

X11 - Operating Expense Rate: Operating Expenses/Net Sales

X12 - Research and development expense rate: (Research and Development Expenses)/Net Sales

X13 - Cash flow rate: Cash Flow from Operating/Current Liabilities

X14 - Interest-bearing debt interest rate: Interest-bearing Debt/Equity

X15 - Tax rate (A): Effective Tax Rate

X16 - Net Value Per Share (B): Book Value Per Share(B)

X17 - Net Value Per Share (A): Book Value Per Share(A)
 X18 - Net Value Per Share (C): Book Value Per Share(C)
 X19 - Persistent EPS in the Last Four Seasons: EPS-Net Income
 X20 - Cash Flow Per Share
 X21 - Revenue Per Share (Yuan ¥): Sales Per Share
 X22 - Operating Profit Per Share (Yuan ¥): Operating Income Per Share
 X23 - Per Share Net profit before tax (Yuan ¥): Pretax Income Per Share
 X24 - Realized Sales Gross Profit Growth Rate
 X25 - Operating Profit Growth Rate: Operating Income Growth
 X26 - After-tax Net Profit Growth Rate: Net Income Growth
 X27 - Regular Net Profit Growth Rate: Continuing Operating Income after Tax Growth
 X28 - Continuous Net Profit Growth Rate: Net Income-Excluding Disposal Gain or Loss Growth
 X29 - Total Asset Growth Rate: Total Asset Growth
 X30 - Net Value Growth Rate: Total Equity Growth
 X31 - Total Asset Return Growth Rate Ratio: Return on Total Asset Growth
 X32 - Cash Reinvestment %: Cash Reinvestment Ratio
 X33 - Current Ratio
 X34 - Quick Ratio: Acid Test
 X35 - Interest Expense Ratio: Interest Expenses/Total Revenue
 X36 - Total debt/Total net worth: Total Liability/Equity Ratio
 X37 - Debt ratio %: Liability/Total Assets
 X38 - Net worth/Assets: Equity/Total Assets
 X39 - Long-term fund suitability ratio (A): (Long-term Liability+Equity)/Fixed Assets
 X40 - Borrowing dependency: Cost of Interest-bearing Debt
 X41 - Contingent liabilities/Net worth: Contingent Liability/Equity
 X42 - Operating profit/Paid-in capital: Operating Income/Capital
 X43 - Net profit before tax/Paid-in capital: Pretax Income/Capital
 X44 - Inventory and accounts receivable/Net value: (Inventory+Accounts Receivables)/Equity
 X45 - Total Asset Turnover
 X46 - Accounts Receivable Turnover
 X47 - Average Collection Days: Days Receivable Outstanding
 X48 - Inventory Turnover Rate (times)
 X49 - Fixed Assets Turnover Frequency
 X50 - Net Worth Turnover Rate (times): Equity Turnover

X51 - Revenue per person: Sales Per Employee
 X52 - Operating profit per person: Operation Income Per Employee
 X53 - Allocation rate per person: Fixed Assets Per Employee
 X54 - Working Capital to Total Assets
 X55 - Quick Assets/Total Assets
 X56 - Current Assets/Total Assets
 X57 - Cash/Total Assets
 X58 - Quick Assets/Current Liability
 X59 - Cash/Current Liability
 X60 - Current Liability to Assets
 X61 - Operating Funds to Liability
 X62 - Inventory/Working Capital
 X63 - Inventory/Current Liability
 X64 - Current Liabilities/Liability
 X65 - Working Capital/Equity
 X66 - Current Liabilities/Equity
 X67 - Long-term Liability to Current Assets
 X68 - Retained Earnings to Total Assets
 X69 - Total income/Total expense
 X70 - Total expense/Assets
 X71 - Current Asset Turnover Rate: Current Assets to Sales
 X72 - Quick Asset Turnover Rate: Quick Assets to Sales
 X73 - Working capital Turnover Rate: Working Capital to Sales
 X74 - Cash Turnover Rate: Cash to Sales
 X75 - Cash Flow to Sales
 X76 - Fixed Assets to Assets
 X77 - Current Liability to Liability
 X78 - Current Liability to Equity
 X79 - Equity to Long-term Liability
 X80 - Cash Flow to Total Assets
 X81 - Cash Flow to Liability
 X82 - CFO to Assets
 X83 - Cash Flow to Equity
 X84 - Current Liability to Current Assets
 X85 - Liability-Assets Flag: 1 if Total Liability exceeds Total Assets, 0 otherwise
 X86 - Net Income to Total Assets
 X87 - Total assets to GNP price
 X88 - No-credit Interval
 X89 - Gross Profit to Sales
 X90 - Net Income to Stockholder's Equity
 X91 - Liability to Equity
 X92 - Degree of Financial Leverage (DFL)
 X93 - Interest Coverage Ratio (Interest expense to EBIT)
 X94 - Net Income Flag: 1 if Net Income is Negative for the last two years, 0 otherwise
 X95 - Equity to Liability

```
[89]: # Import libraries here
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns

import ipywidgets as widgets
from ipywidgets import interact

from imblearn.over_sampling import RandomOverSampler
from imblearn.under_sampling import RandomUnderSampler

from sklearn.impute import SimpleImputer
from sklearn.metrics import (
    ConfusionMatrixDisplay,
    classification_report,
    confusion_matrix,
)
from sklearn.model_selection import GridSearchCV, \
    cross_val_score, train_test_split
from sklearn.pipeline import make_pipeline
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
```

```
[3]: df=pd.read_csv('data/data.csv')
df.head()
```

```
[3]: Bankrupt?    ROA(C) before interest and depreciation before interest \
0             1             0.370594
1             1             0.464291
2             1             0.426071
3             1             0.399844
4             1             0.465022

ROA(A) before interest and % after tax \
0             0.424389
1             0.538214
2             0.499019
3             0.451265
4             0.538432

ROA(B) before interest and depreciation after tax \
0             0.405750
1             0.516730
2             0.472295
3             0.457733
4             0.522298
```

	Operating Gross Margin	Realized Sales Gross Margin \
0	0.601457	0.601457
1	0.610235	0.610235
2	0.601450	0.601364
3	0.583541	0.583541
4	0.598783	0.598783

	Operating Profit Rate	Pre-tax net Interest Rate \
0	0.998969	0.796887
1	0.998946	0.797380
2	0.998857	0.796403
3	0.998700	0.796967
4	0.998973	0.797366

	After-tax net Interest Rate	Non-industry income and expenditure/revenue \
0	0.808809	0.302646
1	0.809301	0.303556
2	0.808388	0.302035
3	0.808966	0.303350
4	0.809304	0.303475

...	Net Income to Total Assets	Total assets to GNP price \
0	0.716845	0.009219
1	0.795297	0.008323
2	0.774670	0.040003
3	0.739555	0.003252
4	0.795016	0.003878

	No-credit Interval	Gross Profit to Sales \
0	0.622879	0.601453
1	0.623652	0.610237
2	0.623841	0.601449
3	0.622929	0.583538
4	0.623521	0.598782

	Net Income to Stockholder's Equity	Liability to Equity \
0	0.827890	0.290202
1	0.839969	0.283846
2	0.836774	0.290189
3	0.834697	0.281721
4	0.839973	0.278514

	Degree of Financial Leverage (DFL) \
0	0.026601
1	0.264577
2	0.026555
3	0.026697

```

4                                0.024752

      Interest Coverage Ratio (Interest expense to EBIT)  Net Income Flag \
0                                0.564050                                1
1                                0.570175                                1
2                                0.563706                                1
3                                0.564663                                1
4                                0.575617                                1

      Equity to Liability
0                                0.016469
1                                0.020794
2                                0.016474
3                                0.023982
4                                0.035490

[5 rows x 96 columns]

```

3 Explore

```

[4]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6819 entries, 0 to 6818
Data columns (total 96 columns):
 #   Column                                Non-Null Count
Dtype  -----
-----
0    Bankrupt?                            6819 non-null
int64
1    ROA(C) before interest and depreciation before interest  6819 non-null
float64
2    ROA(A) before interest and % after tax                    6819 non-null
float64
3    ROA(B) before interest and depreciation after tax         6819 non-null
float64
4    Operating Gross Margin                                    6819 non-null
float64
5    Realized Sales Gross Margin                                6819 non-null
float64
6    Operating Profit Rate                                       6819 non-null
float64
7    Pre-tax net Interest Rate                                   6819 non-null
float64
8    After-tax net Interest Rate                                6819 non-null
float64

```

9	Non-industry income and expenditure/revenue	6819 non-null
float64		
10	Continuous interest rate (after tax)	6819 non-null
float64		
11	Operating Expense Rate	6819 non-null
float64		
12	Research and development expense rate	6819 non-null
float64		
13	Cash flow rate	6819 non-null
float64		
14	Interest-bearing debt interest rate	6819 non-null
float64		
15	Tax rate (A)	6819 non-null
float64		
16	Net Value Per Share (B)	6819 non-null
float64		
17	Net Value Per Share (A)	6819 non-null
float64		
18	Net Value Per Share (C)	6819 non-null
float64		
19	Persistent EPS in the Last Four Seasons	6819 non-null
float64		
20	Cash Flow Per Share	6819 non-null
float64		
21	Revenue Per Share (Yuan ¥)	6819 non-null
float64		
22	Operating Profit Per Share (Yuan ¥)	6819 non-null
float64		
23	Per Share Net profit before tax (Yuan ¥)	6819 non-null
float64		
24	Realized Sales Gross Profit Growth Rate	6819 non-null
float64		
25	Operating Profit Growth Rate	6819 non-null
float64		
26	After-tax Net Profit Growth Rate	6819 non-null
float64		
27	Regular Net Profit Growth Rate	6819 non-null
float64		
28	Continuous Net Profit Growth Rate	6819 non-null
float64		
29	Total Asset Growth Rate	6819 non-null
float64		
30	Net Value Growth Rate	6819 non-null
float64		
31	Total Asset Return Growth Rate Ratio	6819 non-null
float64		
32	Cash Reinvestment %	6819 non-null
float64		

33	Current Ratio	6819 non-null
	float64	
34	Quick Ratio	6819 non-null
	float64	
35	Interest Expense Ratio	6819 non-null
	float64	
36	Total debt/Total net worth	6819 non-null
	float64	
37	Debt ratio %	6819 non-null
	float64	
38	Net worth/Assets	6819 non-null
	float64	
39	Long-term fund suitability ratio (A)	6819 non-null
	float64	
40	Borrowing dependency	6819 non-null
	float64	
41	Contingent liabilities/Net worth	6819 non-null
	float64	
42	Operating profit/Paid-in capital	6819 non-null
	float64	
43	Net profit before tax/Paid-in capital	6819 non-null
	float64	
44	Inventory and accounts receivable/Net value	6819 non-null
	float64	
45	Total Asset Turnover	6819 non-null
	float64	
46	Accounts Receivable Turnover	6819 non-null
	float64	
47	Average Collection Days	6819 non-null
	float64	
48	Inventory Turnover Rate (times)	6819 non-null
	float64	
49	Fixed Assets Turnover Frequency	6819 non-null
	float64	
50	Net Worth Turnover Rate (times)	6819 non-null
	float64	
51	Revenue per person	6819 non-null
	float64	
52	Operating profit per person	6819 non-null
	float64	
53	Allocation rate per person	6819 non-null
	float64	
54	Working Capital to Total Assets	6819 non-null
	float64	
55	Quick Assets/Total Assets	6819 non-null
	float64	
56	Current Assets/Total Assets	6819 non-null
	float64	

57	Cash/Total Assets	6819 non-null
	float64	
58	Quick Assets/Current Liability	6819 non-null
	float64	
59	Cash/Current Liability	6819 non-null
	float64	
60	Current Liability to Assets	6819 non-null
	float64	
61	Operating Funds to Liability	6819 non-null
	float64	
62	Inventory/Working Capital	6819 non-null
	float64	
63	Inventory/Current Liability	6819 non-null
	float64	
64	Current Liabilities/Liability	6819 non-null
	float64	
65	Working Capital/Equity	6819 non-null
	float64	
66	Current Liabilities/Equity	6819 non-null
	float64	
67	Long-term Liability to Current Assets	6819 non-null
	float64	
68	Retained Earnings to Total Assets	6819 non-null
	float64	
69	Total income/Total expense	6819 non-null
	float64	
70	Total expense/Assets	6819 non-null
	float64	
71	Current Asset Turnover Rate	6819 non-null
	float64	
72	Quick Asset Turnover Rate	6819 non-null
	float64	
73	Working capitol Turnover Rate	6819 non-null
	float64	
74	Cash Turnover Rate	6819 non-null
	float64	
75	Cash Flow to Sales	6819 non-null
	float64	
76	Fixed Assets to Assets	6819 non-null
	float64	
77	Current Liability to Liability	6819 non-null
	float64	
78	Current Liability to Equity	6819 non-null
	float64	
79	Equity to Long-term Liability	6819 non-null
	float64	
80	Cash Flow to Total Assets	6819 non-null
	float64	

```


    81  Cash Flow to Liability                                6819 non-null
float64
    82  CFO to Assets                                        6819 non-null
float64
    83  Cash Flow to Equity                                6819 non-null
float64
    84  Current Liability to Current Assets                 6819 non-null
float64
    85  Liability-Assets Flag                               6819 non-null
int64
    86  Net Income to Total Assets                         6819 non-null
float64
    87  Total assets to GNP price                           6819 non-null
float64
    88  No-credit Interval                                 6819 non-null
float64
    89  Gross Profit to Sales                              6819 non-null
float64
    90  Net Income to Stockholder's Equity                 6819 non-null
float64
    91  Liability to Equity                                6819 non-null
float64
    92  Degree of Financial Leverage (DFL)                 6819 non-null
float64
    93  Interest Coverage Ratio (Interest expense to EBIT) 6819 non-null
float64
    94  Net Income Flag                                    6819 non-null
int64
    95  Equity to Liability                                6819 non-null
float64
dtypes: float64(93), int64(3)
memory usage: 5.0 MB

```

That's solid information. We know all our features are numerical and that we have no missing data. But, it's a good idea to do some visualizations to see if there are any interesting trends or ideas we should keep in mind while we work. First, let's take a look at how many firms are bankrupt, and how many are not.

```
[5]: df['Bankrupt?'].value_counts(normalize=True).plot(kind='bar')
```

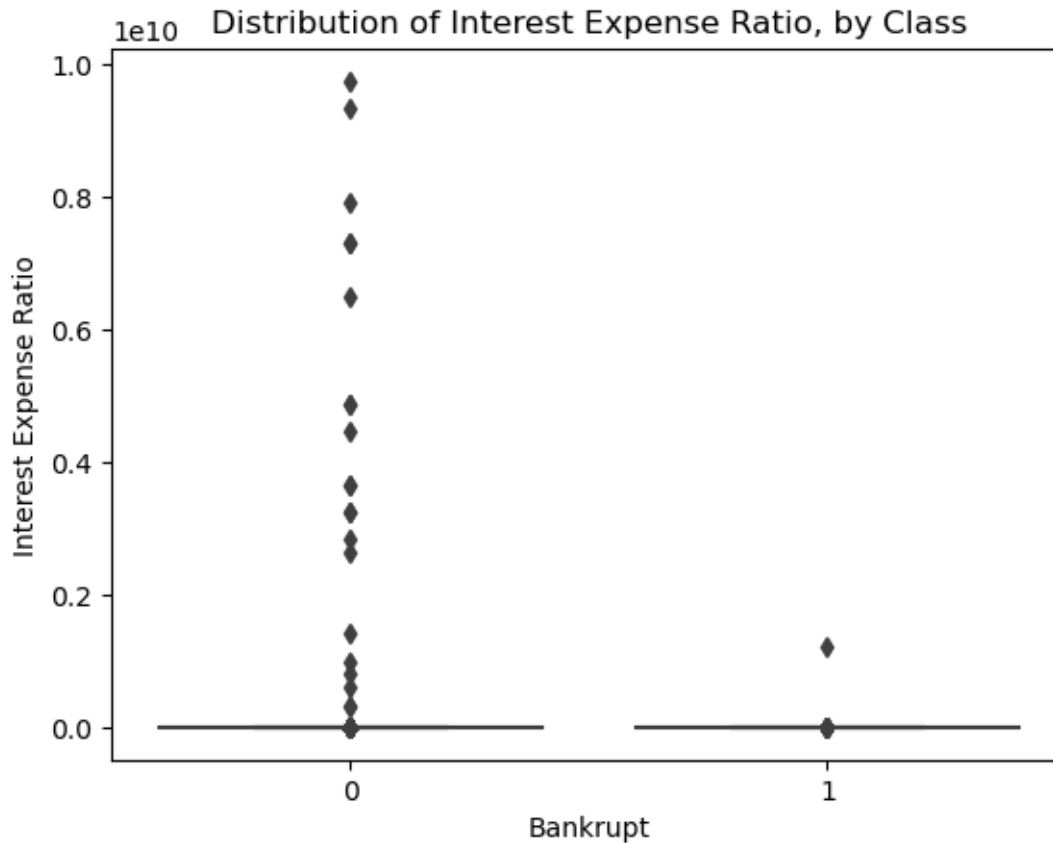
```
[5]: <AxesSubplot:>
```



That's good news for Japan's economy! Since it looks like most of the companies in our dataset are doing all right for themselves, let's drill down a little farther. However, it also shows us that we have an **imbalanced** dataset, where our majority class is far bigger than our minority class.

there were 64 features of each company, each of which had some kind of numerical value. It might be useful to understand where the values for one of these features cluster, so let's make a boxplot to see how the values in "Interest Expense Ratio" are distributed.

```
[6]: # Create boxplot
sns.boxplot(y='Accounts Receivable Turnover',x='Bankrupt?',data=df)
plt.xlabel("Bankrupt")
plt.ylabel("Interest Expense Ratio")
plt.title("Distribution of Interest Expense Ratio, by Class");
```




Why does this look so funny? Remember that boxplots exist to help us see the quartiles in a dataset, and this one doesn't really do that. Let's check the distribution of "Interest Expense Ratio" to see if we can figure out what's going on here.

```
[7]: # Summary statistics for `Operating Profit Rate`
df[' Operating Profit Rate'].describe().apply('{0:,.0f}'.format)
```

```
[7]: count      6,819
      mean         1
      std          0
      min          0
      25%          1
      50%          1
      75%          1
      max          1
      Name: Operating Profit Rate, dtype: object
```

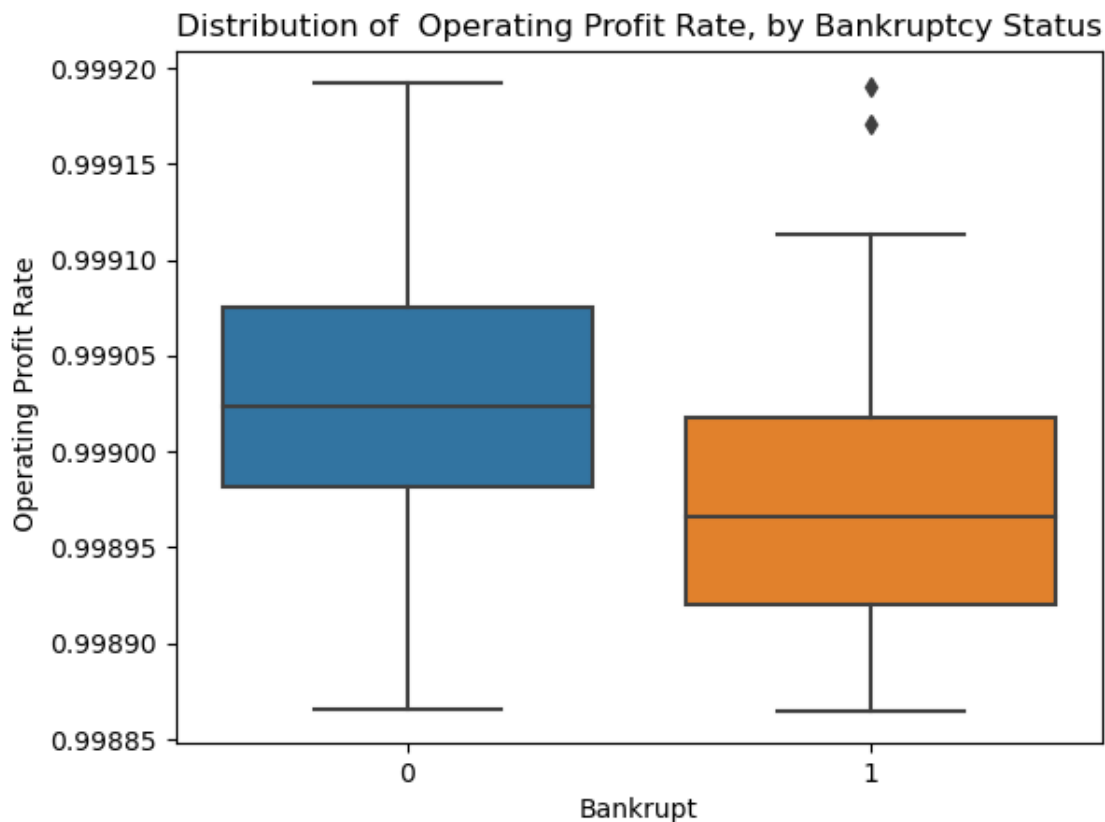
```
[8]: df[' Operating Profit Rate'].hist()
```

```
[8]: <AxesSubplot:>
```



Aha! We saw it in the numbers and now we see it in the histogram. The data is very skewed. So, in order to create a helpful boxplot, we need to trim the data.

```
[9]: # Create clipped boxplot
q1,q9=df[' Operating Profit Rate'].quantile([0.1,0.9])
mask=df[' Operating Profit Rate'].between(q1,q9)
sns.boxplot(x='Bankrupt?',y=' Operating Profit Rate',data=df[mask])
plt.xlabel("Bankrupt")
plt.ylabel(" Operating Profit Rate")
plt.title("Distribution of Operating Profit Rate, by Bankruptcy Status");
```

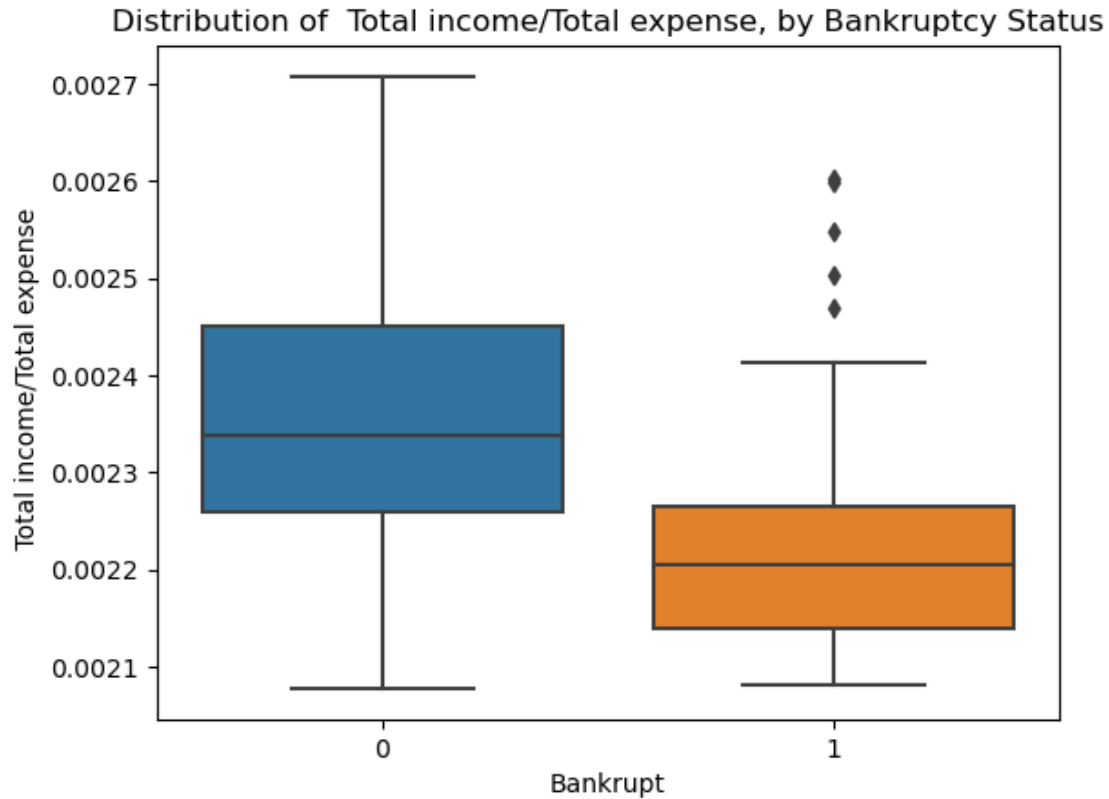


That makes a lot more sense. Let's take a look at some of the other features in the dataset to see what else is out there.

```
[10]: df[' Total income/Total expense'].describe().apply('{0:,.0f}'.format)
```

```
[10]: count      6,819
      mean         0
      std         0
      min         0
      25%         0
      50%         0
      75%         0
      max         1
      Name: Total income/Total expense, dtype: object
```

```
[11]: q1,q9=df[' Total income/Total expense'].quantile([0.1,0.9])
      mask=df[' Total income/Total expense'].between(q1,q9)
      sns.boxplot(x='Bankrupt?',y=' Total income/Total expense',data=df[mask])
      plt.xlabel("Bankrupt")
      plt.ylabel(" Total income/Total expense")
      plt.title("Distribution of Total income/Total expense, by Bankruptcy Status");
```

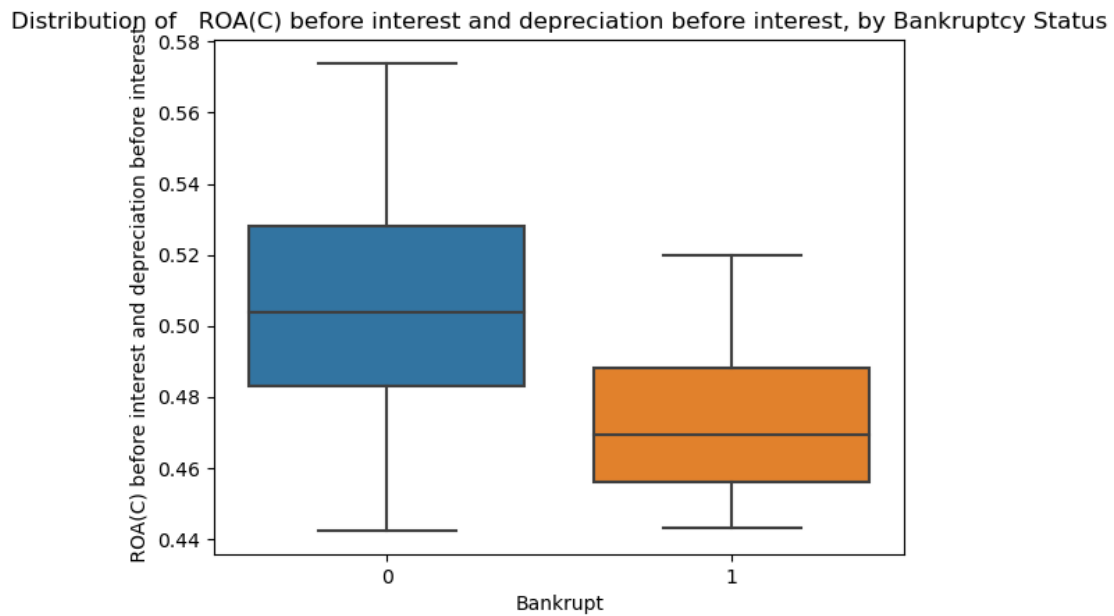


```
[12]: df[' ROA(C) before interest and depreciation before interest'].describe().
      ↪ apply('{0:,.0f}'.format)
```

```
[12]: count      6,819
      mean         1
      std          0
      min          0
      25%          0
      50%          1
      75%          1
      max          1
      Name: ROA(C) before interest and depreciation before interest, dtype: object
```

```
[13]: q1,q9=df[' ROA(C) before interest and depreciation before interest'].
      ↪ quantile([0.1,0.9])
      mask=df[' ROA(C) before interest and depreciation before interest'].
      ↪ between(q1,q9)
      sns.boxplot(x='Bankrupt?',y=' ROA(C) before interest and depreciation before_
      ↪ interest',data=df[mask])
      plt.xlabel("Bankrupt")
      plt.ylabel(" ROA(C) before interest and depreciation before interest")
```

```
plt.title("Distribution of ROA(C) before interest and depreciation before interest, by Bankruptcy Status");
```

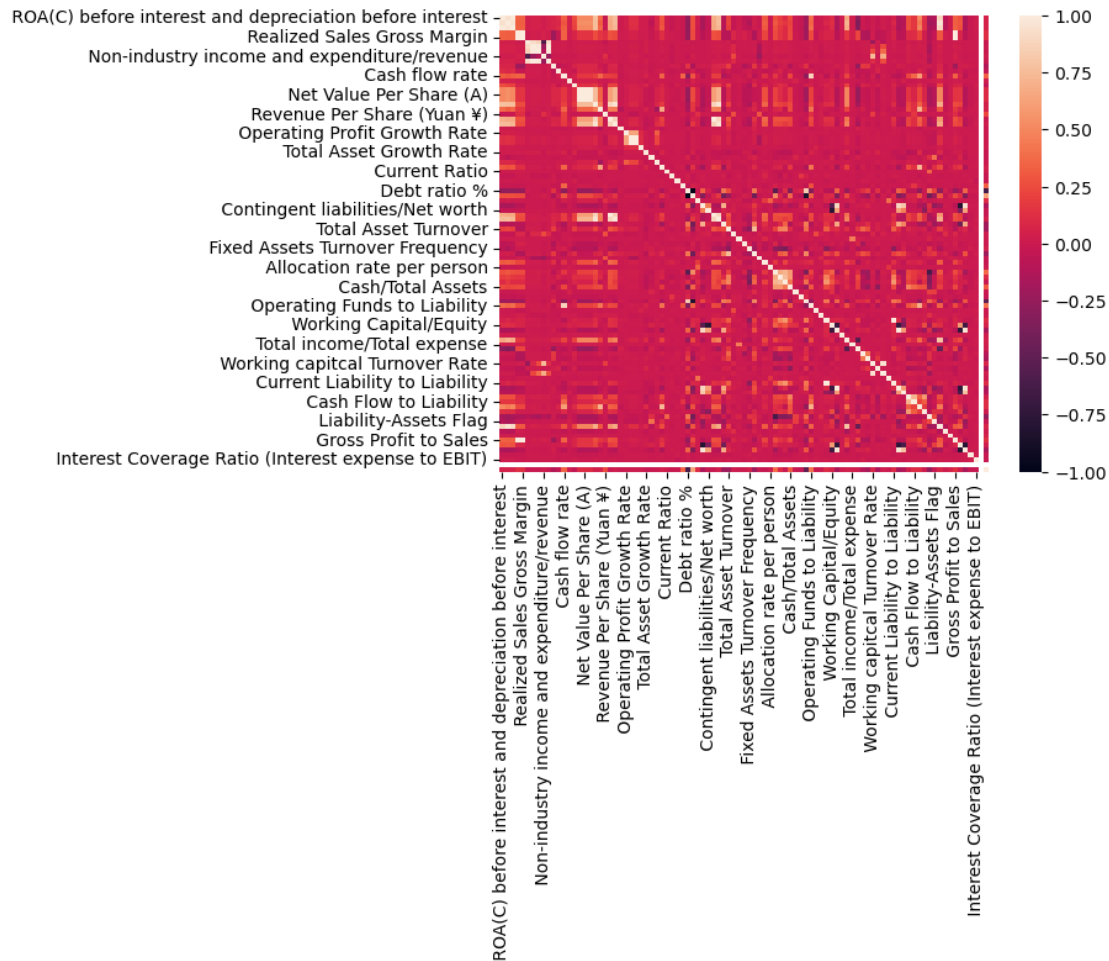


Looking at other features, we can see that they're skewed, too. This will be important to keep in mind when we decide what type of model we want to use.

Another important consideration for model selection is whether there are any issues with multicollinearity in our model. Let's check.

```
[14]: corr=df.drop(columns='Bankrupt?').corr()
      sns.heatmap(corr)
```

```
[14]: <AxesSubplot:>
```

So what did we learn from this EDA? First, our data is imbalanced. This is something we need to address in our data preparation. Second, our features haven't missing but if there are values that we'll need to impute. And since the features are highly skewed, the best imputation strategy is likely median, not mean. Finally, we have autocorrelation issues, which means that we should steer clear of linear models, and try a tree-based model instead.

4 Split

```
[15]: target = "Bankrupt?"
X = df.drop(columns=target)
y = df[target]

print("X shape:", X.shape)
print("y shape:", y.shape)
```

```
X shape: (6819, 95)
y shape: (6819,)
```

```
[16]: X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.
      ↪2,random_state=42)

print("X_train shape:", X_train.shape)
print("y_train shape:", y_train.shape)
print("X_test shape:", X_test.shape)
print("y_test shape:", y_test.shape)
```

```
X_train shape: (5455, 95)
y_train shape: (5455,)
X_test shape: (1364, 95)
y_test shape: (1364,)
```

5 Resample

Now that we've split our data into training and validation sets, we can address the class imbalance we saw during our EDA. One strategy is to resample the training data. There are many to do this, so let's start with under-sampling.

```
[17]: under_sampler =RandomUnderSampler(random_state=42)
      X_train_under, y_train_under = under_sampler.fit_resample(X_train,y_train)
      print(X_train_under.shape)
      X_train_under.head()
```

```
(338, 95)
```

```
[17]:      ROA(C) before interest and depreciation before interest \
5782                                     0.547409
5200                                     0.575196
3786                                     0.490567
1338                                     0.543899
4804                                     0.608980
```

```
      ROA(A) before interest and % after tax \
5782                                     0.597580
5200                                     0.622928
3786                                     0.552933
1338                                     0.598779
4804                                     0.658689
```

```
      ROA(B) before interest and depreciation after tax \
5782                                     0.583757
5200                                     0.609187
3786                                     0.539483
1338                                     0.593180
4804                                     0.655174
```

```
      Operating Gross Margin   Realized Sales Gross Margin \
```

5782	0.604707	0.604707
5200	0.606142	0.606142
3786	0.629665	0.629520
1338	0.602740	0.602754
4804	0.612491	0.613183

	Operating Profit Rate	Pre-tax net Interest Rate \
5782	0.999089	0.797593
5200	0.999085	0.797859
3786	0.999108	0.797607
1338	0.999041	0.797495
4804	0.999193	0.798374

	After-tax net Interest Rate \
5782	0.809456
5200	0.809654
3786	0.809484
1338	0.809399
4804	0.810111

	Non-industry income and expenditure/revenue \
5782	0.303629
5200	0.304102
3786	0.303613
1338	0.303559
4804	0.304774

	Continuous interest rate (after tax) ...	Net Income to Total Assets \
5782	0.781726 ...	0.831586
5200	0.781934 ...	0.845096
3786	0.781755 ...	0.810341
1338	0.781664 ...	0.828565
4804	0.782440 ...	0.867959

	Total assets to GNP price	No-credit Interval	Gross Profit to Sales \
5782	0.001479	0.624287	0.604706
5200	0.001891	0.624553	0.606139
3786	0.001115	0.624214	0.629664
1338	0.001621	0.623379	0.602735
4804	0.010567	0.621156	0.612491

	Net Income to Stockholder's Equity	Liability to Equity \
5782	0.842928	0.279743
5200	0.843117	0.277089
3786	0.840741	0.275084
1338	0.843240	0.282201
4804	0.843773	0.275537

	Degree of Financial Leverage (DFL) \
5782	0.026801
5200	0.026791
3786	0.026791
1338	0.026865
4804	0.026791

	Interest Coverage Ratio (Interest expense to EBIT)	Net Income Flag \
5782	0.565205	1
5200	0.565159	1
3786	0.565158	1
1338	0.565489	1
4804	0.565158	1

	Equity to Liability
5782	0.029354
5200	0.050445
3786	0.250781
1338	0.023104
4804	0.123342

[5 rows x 95 columns]

And then we'll over-sample.

```
[18]: over_sampler = RandomOverSampler(random_state=42)
X_train_over, y_train_over = over_sampler.fit_resample(X_train,y_train)
print(X_train_over.shape)
X_train_over.head()
```

(10572, 95)

```
[18]: ROA(C) before interest and depreciation before interest \
0      0.498513
1      0.506606
2      0.508799
3      0.499976
4      0.477892
```

```
ROA(A) before interest and % after tax \
0      0.542848
1      0.562309
2      0.561001
3      0.562527
4      0.547700
```

```
ROA(B) before interest and depreciation after tax \
```

0	0.544622
1	0.558863
2	0.554687
3	0.546764
4	0.529150

	Operating Gross Margin	Realized Sales Gross Margin \
0	0.599194	0.599036
1	0.609334	0.609334
2	0.614242	0.614055
3	0.597825	0.597825
4	0.600362	0.600362

	Operating Profit Rate	Pre-tax net Interest Rate \
0	0.998986	0.797412
1	0.999027	0.797450
2	0.999094	0.797533
3	0.999004	0.797411
4	0.998975	0.797412

	After-tax net Interest Rate	Non-industry income and expenditure/revenue \
0	0.809330	0.303528
1	0.809375	0.303508
2	0.809424	0.303514
3	0.809329	0.303490
4	0.809333	0.303551

	Continuous interest rate (after tax) ...	Net Income to Total Assets \
0	0.781593 ...	0.801313
1	0.781637 ...	0.810914
2	0.781692 ...	0.809740
3	0.781590 ...	0.810082
4	0.781584 ...	0.804638

	Total assets to GNP price	No-credit Interval	Gross Profit to Sales \
0	0.005821	0.623649	0.599196
1	0.000481	0.623932	0.609332
2	0.001397	0.623714	0.614241
3	0.000998	0.623986	0.597824
4	0.002826	0.623845	0.600363

	Net Income to Stockholder's Equity	Liability to Equity \
0	0.840580	0.282564
1	0.841339	0.280570
2	0.840969	0.277772
3	0.841885	0.286871
4	0.840885	0.282073

	Degree of Financial Leverage (DFL) \	
0	0.027239	
1	0.026843	
2	0.026864	
3	0.026951	
4	0.026959	

	Interest Coverage Ratio (Interest expense to EBIT)	Net Income Flag \
0	0.566658	1
1	0.565395	1
2	0.565484	1
3	0.565820	1
4	0.565848	1

	Equity to Liability
0	0.022512
1	0.026670
2	0.041556
3	0.018173
4	0.023328

[5 rows x 95 columns]

6 Build Model

6.1 base line

```
[19]: acc_baseline = y_train.value_counts(normalize=True).max()
print("Baseline Accuracy:", round(acc_baseline, 4))
```

Baseline Accuracy: 0.969

Note here that, because our classes are imbalanced, the baseline accuracy is very high. We should keep this in mind because, even if our trained model gets a high validation accuracy score, that doesn't mean it's actually *good*.

6.2 Iterate

6.3 1- Decision Tree model

```
[20]: # Fit on `X_train`, `y_train`
model_reg =
    ↳make_pipeline(SimpleImputer(strategy='median'),DecisionTreeClassifier(random_state=42))
model_reg.fit(X_train, y_train)

# Fit on `X_train_under`, `y_train_under`
```

```

model_under =
    ↳make_pipeline(SimpleImputer(strategy='median'),DecisionTreeClassifier(random_state=42))
model_under.fit(X_train_under, y_train_under)

# Fit on `X_train_over`, `y_train_over`
model_over =
    ↳make_pipeline(SimpleImputer(strategy='median'),DecisionTreeClassifier(random_state=42))
model_over.fit(X_train_over, y_train_over)

```

```

[20]: Pipeline(steps=[('simpleimputer', SimpleImputer(strategy='median')),
                        ('decisiontreeclassifier',
                         DecisionTreeClassifier(random_state=42))])

```

6.4 Evaluate

```

[21]: for m in [model_reg, model_under, model_over]:
        acc_train = m.score(X_train,y_train)
        acc_test = m.score(X_test,y_test)

        print("Training Accuracy:", round(acc_train, 4))
        print("Test Accuracy:", round(acc_test, 4))

```

```

Training Accuracy: 1.0
Test Accuracy: 0.9531
Training Accuracy: 0.813
Test Accuracy: 0.8057
Training Accuracy: 1.0
Test Accuracy: 0.9604

```

“good” accuracy scores don’t tell us much about the model’s performance when dealing with imbalanced data. So instead of looking at what the model got right or wrong

```


[22]: # Plot confusion matrix
ConfusionMatrixDisplay.from_estimator(model_reg,X_test,y_test)

```

```

[22]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x19f73b43220>

```



```
[23]: #Determine the depth of the decision tree in model_over
depth = model_over.named_steps['decisiontreeclassifier'].get_depth()
print(depth)
```

53


6.5 Communicate

Now that we have a reasonable model, let's graph the importance of each feature.

```
[24]: # Get importances
importances = model_over.named_steps['decisiontreeclassifier'].
    feature_importances_

# Put importances into a Series
feat_imp = pd.Series(importances, index=X_train_over.columns).sort_values()

# Plot series
feat_imp.tail(10).plot(kind='barh')
plt.xlabel("Gini Importance")
plt.ylabel("Feature")
plt.title("model_over Feature Importance");
```

confusion matrix does not give us the best result so we need another model to get more *TP*, *TN* results

6.6 2- Random Forest Classifier

```
[26]: clf = make_pipeline(SimpleImputer(), RandomForestClassifier(random_state=42))
      print(clf)
```

```
Pipeline(steps=[('simpleimputer', SimpleImputer()),
                  ('randomforestclassifier',
                   RandomForestClassifier(random_state=42))])
```

```
[27]: cv_acc_scores = cross_val_score(clf, X_train_over, y_train_over, cv=5, n_jobs=-1)
      print(cv_acc_scores)
```

```
[0.99338061 0.99432624 0.99432356 0.99668874 0.99432356]
```

```
[28]: params = {'simpleimputer__strategy': ['mean', 'median'],
                'randomforestclassifier__max_depth': range(10, 50, 10),
                'randomforestclassifier__n_estimators': range(25, 100, 25)}
```

```
[29]: model = GridSearchCV(clf, param_grid=params, cv=5, n_jobs=-1, verbose=1)
      model.fit(X_train_over, y_train_over)
```

Fitting 5 folds for each of 24 candidates, totalling 120 fits

```
[29]: GridSearchCV(cv=5,
                  estimator=Pipeline(steps=[('simpleimputer', SimpleImputer()),
                                              ('randomforestclassifier',
                                               RandomForestClassifier(random_state=42))]),
```

```

n_jobs=-1,
param_grid={'randomforestclassifier__max_depth': range(10, 50, 10),
            'randomforestclassifier__n_estimators': range(25, 100,
25),
            'simpleimputer__strategy': ['mean', 'median']},
verbose=1)

```

```

[32]: cv_results = pd.DataFrame(model.cv_results_)
      cv_results.head(5)

```

```

[32]:  mean_fit_time  std_fit_time  mean_score_time  std_score_time  \
0      1.287459      0.049826      0.034180      0.003967
1      1.458360      0.024056      0.035780      0.005841
2      2.503959      0.076076      0.069960      0.025022
3      3.033654      0.096780      0.054369      0.011032
4      3.993904      0.078464      0.079954      0.013805

      param_randomforestclassifier__max_depth  \
0                                           10
1                                           10
2                                           10
3                                           10
4                                           10

      param_randomforestclassifier__n_estimators  param_simpleimputer__strategy  \
0                                           25                                mean
1                                           25                                median
2                                           50                                mean
3                                           50                                median
4                                           75                                mean

      params  split0_test_score  \
0  {'randomforestclassifier__max_depth': 10, 'ran...      0.979196
1  {'randomforestclassifier__max_depth': 10, 'ran...      0.979196
2  {'randomforestclassifier__max_depth': 10, 'ran...      0.979669
3  {'randomforestclassifier__max_depth': 10, 'ran...      0.979669
4  {'randomforestclassifier__max_depth': 10, 'ran...      0.979196

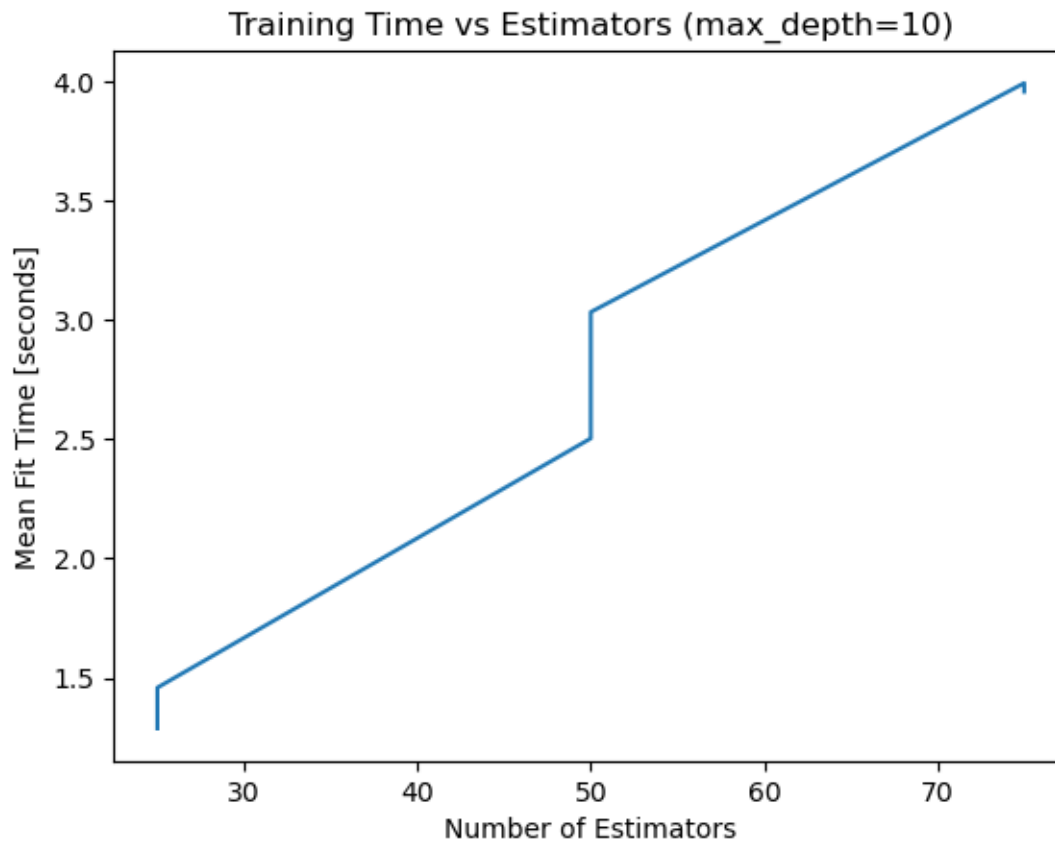
      split1_test_score  split2_test_score  split3_test_score  split4_test_score  \
0      0.977778      0.980132      0.978713      0.979659
1      0.977778      0.980132      0.978713      0.979659
2      0.979196      0.980132      0.979186      0.979186
3      0.979196      0.980132      0.979186      0.979186
4      0.980142      0.977294      0.979659      0.978713

      mean_test_score  std_test_score  rank_test_score
0      0.979096      0.000811      21

```

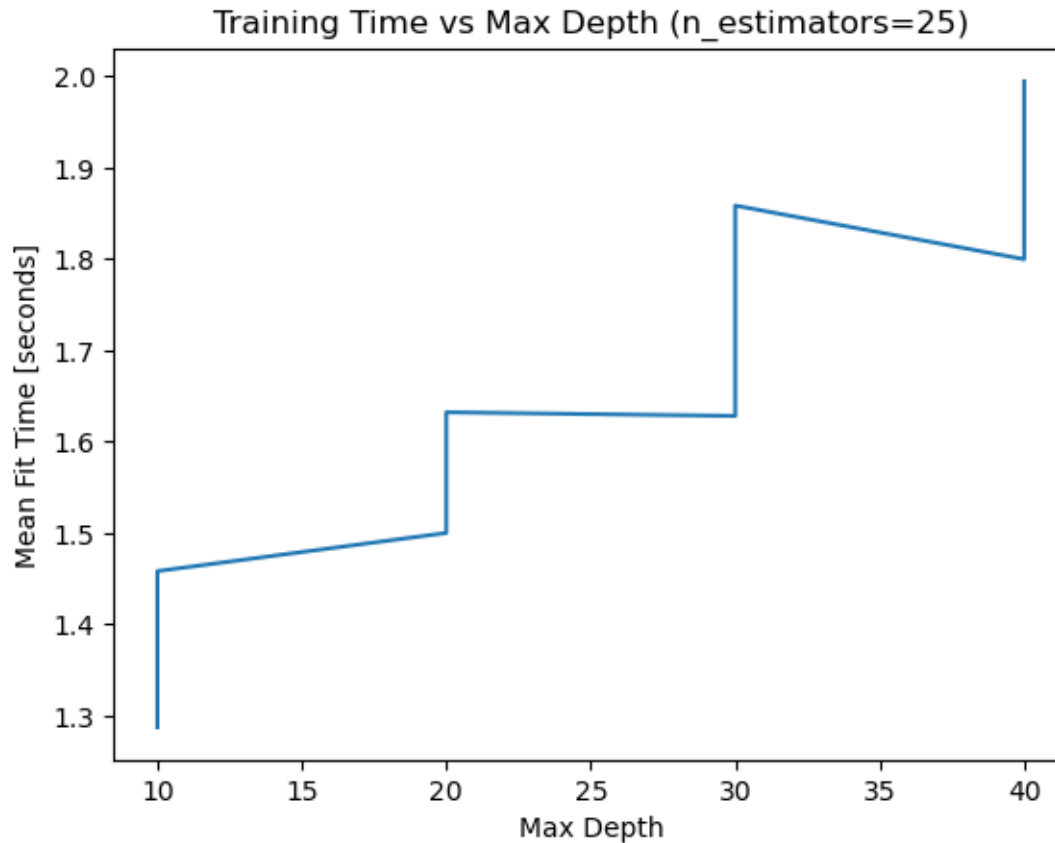
1	0.979096	0.000811	21
2	0.979474	0.000378	19
3	0.979474	0.000378	19
4	0.979001	0.000977	23

```
[36]: # Create mask
mask = cv_results['param_randomforestclassifier__max_depth']==10
# Plot fit time vs n_estimators
plt.
    plot(cv_results[mask]['param_randomforestclassifier__n_estimators'],cv_results[mask]['mean_
# Label axes
plt.xlabel("Number of Estimators")
plt.ylabel("Mean Fit Time [seconds]")
plt.title("Training Time vs Estimators (max_depth=10)");
```



```
[38]: # Create mask
mask = cv_results['param_randomforestclassifier__n_estimators']==25
# Plot fit time vs max_depth
plt.
    plot(cv_results[mask]['param_randomforestclassifier__max_depth'],cv_results[mask]['mean_fit
```

```
# Label axes
plt.xlabel("Max Depth")
plt.ylabel("Mean Fit Time [seconds]")
plt.title("Training Time vs Max Depth (n_estimators=25)");
```



There's a general upwards trend, but we see a lot of up-and-down here. That's because for each max depth, grid search tries two different imputation strategies: mean and median. Median is a lot faster to calculate, so that speeds up training time.

Finally, let's look at the hyperparameters that led to the best performance.

```
[39]: # Extract best hyperparameters
model.best_params_
```

```
[39]: {'randomforestclassifier__max_depth': 40,
      'randomforestclassifier__n_estimators': 50,
      'simpleimputer__strategy': 'mean'}
```

6.7 Evaluate

```
[40]: acc_train = model.score(X_train,y_train)
      acc_test = model.score(X_test,y_test)


      print("Training Accuracy:", round(acc_train, 4))
      print("Test Accuracy:", round(acc_test, 4))
```

Training Accuracy: 1.0

Test Accuracy: 0.9663

```
[41]: # Plot confusion matrix
      ConfusionMatrixDisplay.from_estimator(model,X_test,y_test)
```

```
[41]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x19f7395e910>
```




6.8 Communicate

```
[43]: # Get feature names from training data
      features = X_train_over.columns
      # Extract importances from model
```

```

importances = model.best_estimator_.named_steps['randomforestclassifier'].
    feature_importances_
# Create a series with feature names and importances
feat_imp = pd.Series(importances,index=features).sort_values()
# Plot 10 most important features
feat_imp.tail(10).plot(kind='barh')
plt.xlabel("Gini Importance")
plt.ylabel("Feature")
plt.title("Feature Importance");

```



```

[45]: # save the model
import pickle
with open( "Random_Forest_model.pkl",'wb') as f :
    pickle.dump(model,f)

```

```

[65]: X_test.to_csv("data/X_test.csv",index=False)

```

```

[66]: def make_predictions(data_filepath, model_filepath):
    # Wrangle JSON file
    X_test = pd.read_csv(data_filepath)
    # Load model
    with open(model_filepath,'rb') as file:
        model=pickle.load(file)
    # Generate predictions
    y_test_pred = model.predict(X_test)
    # Put predictions into Series with name "bankrupt", and same index as X_test
    y_test_pred = pd.Series(y_test_pred,index=X_test.index,name='bankrupt')
    return y_test_pred

```

```
[67]: y_test_pred = make_predictions(
        data_filepath="data/X_test.csv",
        model_filepath="Random_Forest_model.pkl",
    )

    print("predictions shape:", y_test_pred.shape)
    y_test_pred.head()
```

predictions shape: (1364,)

```
[67]: 0    0
      1    0
      2    0
      3    1
      4    0
      Name: bankrupt, dtype: int64
```

6.9 3- Gradient Boosting Classifier

```
[71]: clf = make_pipeline(SimpleImputer(), GradientBoostingClassifier())
    print(clf)
```

Pipeline(steps=[('simpleimputer', SimpleImputer()),
('gradientboostingclassifier', GradientBoostingClassifier())])

```
[72]: params = {'simpleimputer__strategy': ['mean', 'median']
               , 'gradientboostingclassifier__max_depth': range(2, 5)
               , 'gradientboostingclassifier__n_estimators': range(20, 31, 5)}
    params
```

```
[72]: {'simpleimputer__strategy': ['mean', 'median'],
      'gradientboostingclassifier__max_depth': range(2, 5),
      'gradientboostingclassifier__n_estimators': range(20, 31, 5)}
```

Note that we're trying much smaller numbers of `n_estimators`. This is because GradientBoostingClassifier is slower to train than the RandomForestClassifier. You can try increasing the number of estimators to see if model performance improves, but keep in mind that you could be waiting a long time!

```
[73]: model = GridSearchCV(clf, param_grid=params, cv=5, n_jobs=-1, verbose=1)
```

```
[74]: # Fit model to over-sampled training data
    model.fit(X_train_over, y_train_over)
```

Fitting 5 folds for each of 18 candidates, totalling 90 fits

```
[74]: GridSearchCV(cv=5,
                  estimator=Pipeline(steps=[('simpleimputer', SimpleImputer()),
                                             ('gradientboostingclassifier',
```

```

GradientBoostingClassifier()))],
n_jobs=-1,
param_grid={'gradientboostingclassifier__max_depth': range(2, 5),
            'gradientboostingclassifier__n_estimators': range(20,
31, 5),
            'simpleimputer__strategy': ['mean', 'median']},
verbose=1)

```

```

[75]: results = pd.DataFrame(model.cv_results_)
      results.sort_values("rank_test_score").head(10)

```

```

[75]:   mean_fit_time  std_fit_time  mean_score_time  std_score_time  \
16      10.763213      0.056432      0.019989      0.005997
17      11.110814      0.148439      0.015791      0.002039
15       9.220324      0.104813      0.016591      0.002243
14       9.269073      0.249547      0.015791      0.000399
13       8.057969      0.085673      0.020789      0.007164
12       8.097546      0.107288      0.016591      0.002330
10       9.478034      0.150210      0.018589      0.002153
11       9.734389      0.149561      0.017191      0.001165
9        8.219004      0.069603      0.016591      0.001019
8        8.365986      0.155735      0.021588      0.008208

```

```

      param_gradientboostingclassifier__max_depth  \
16                                                4
17                                                4
15                                                4
14                                                4
13                                                4
12                                                4
10                                                3
11                                                3
9                                                 3
8                                                 3

```

```

      param_gradientboostingclassifier__n_estimators  \
16                                                  30
17                                                  30
15                                                  25
14                                                  25
13                                                  20
12                                                  20
10                                                  30
11                                                  30
9                                                   25
8                                                   25

```


	param_simpleimputer__strategy \
16	mean
17	median
15	median
14	mean
13	median
12	mean
10	mean
11	median
9	median
8	mean

	params	split0_test_score \
16	{'gradientboostingclassifier__max_depth': 4, '...	0.963121
17	{'gradientboostingclassifier__max_depth': 4, '...	0.962648
15	{'gradientboostingclassifier__max_depth': 4, '...	0.961229
14	{'gradientboostingclassifier__max_depth': 4, '...	0.961229
13	{'gradientboostingclassifier__max_depth': 4, '...	0.957447
12	{'gradientboostingclassifier__max_depth': 4, '...	0.957447
10	{'gradientboostingclassifier__max_depth': 3, '...	0.945154
11	{'gradientboostingclassifier__max_depth': 3, '...	0.945154
9	{'gradientboostingclassifier__max_depth': 3, '...	0.939007
8	{'gradientboostingclassifier__max_depth': 3, '...	0.939007

	split1_test_score	split2_test_score	split3_test_score \
16	0.964066	0.964522	0.966887
17	0.964066	0.964522	0.966414
15	0.962175	0.962157	0.964995
14	0.962175	0.962157	0.964995
13	0.960284	0.957900	0.960738
12	0.960284	0.957900	0.960738
10	0.948936	0.950804	0.947020
11	0.948936	0.950804	0.946547
9	0.942317	0.932829	0.942763
8	0.942317	0.932829	0.942763

	split4_test_score	mean_test_score	std_test_score	rank_test_score
16	0.971618	0.966043	0.003051	1
17	0.972091	0.965948	0.003299	2
15	0.970199	0.964151	0.003277	3
14	0.969726	0.964056	0.003104	4
13	0.966887	0.960651	0.003373	5
12	0.966887	0.960651	0.003373	5
10	0.945601	0.947503	0.002113	7
11	0.945601	0.947408	0.002143	8
9	0.942763	0.939936	0.003822	9
8	0.942763	0.939936	0.003822	9

```
[76]: # Extract best hyperparameters
model.best_params_
```

```
[76]: {'gradientboostingclassifier__max_depth': 4,
      'gradientboostingclassifier__n_estimators': 30,
      'simpleimputer__strategy': 'mean'}
```

6.10 Evaluate


```
[77]: acc_train = model.score(X_train,y_train)
      acc_test = model.score(X_test,y_test)

      print("Training Accuracy:", round(acc_train, 4))
      print("Validation Accuracy:", round(acc_test, 4))
```

Training Accuracy: 0.9487
Validation Accuracy: 0.9267

```
[78]: # Plot confusion matrix
ConfusionMatrixDisplay.from_estimator(model,X_test,y_test)
```

```
[78]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x19f7b4e8880>
```



This matrix is a great reminder of how imbalanced our data is, and of why accuracy isn't always the best metric for judging whether or not a model is giving us what we want. After all, if 95% of the companies in our dataset didn't go bankrupt, all the model has to do is always predict {"bankrupt": 0}, and it'll be right 95% of the time. The accuracy score will be amazing, but it won't tell us what we really need to know.

Instead, we can evaluate our model using two new metrics: precision and recall. The precision score is important when we want our model to only predict that a company will go bankrupt if its very confident in its prediction. The recall score is important if we want to make sure to identify all the companies that will go bankrupt, even if that means being incorrect sometimes.

```
[81]: # Print classification report
print(classification_report(y_test,model.predict(X_test)))
```

	precision	recall	f1-score	support
0	0.99	0.94	0.96	1313
1	0.29	0.67	0.40	51
accuracy			0.93	1364
macro avg	0.64	0.80	0.68	1364
weighted avg	0.96	0.93	0.94	1364

suppose a manager give me a task that that every time i predict tp i will get a profit for my company = 100_000_000 for each tp

and

every time i predict fp i will make a lose for my company = 2500_000_000 for each fp

```
[93]: def make_cnf_matrix(threshold):
    y_pred_prob=model.predict_proba(X_test)[:,-1]
    y_pred=y_pred_prob>threshold
    con=confusion_matrix(y_test,y_pred)
    tn,fp,fn,tp=con.ravel()
    print(f'profit: ${tp*100_000_000}')
    print(f'losses: ${fp*250_000_000}')
    ConfusionMatrixDisplay.from_predictions(y_test,y_pred,colorbar=False)


thresh_widget = widgets.FloatSlider(min=0,max=1,step=0.05,value=0.5)

interact(make_cnf_matrix, threshold=thresh_widget);
```

```
interactive(children=(FloatSlider(value=0.5, description='threshold', max=1.0,
    step=0.05), Output()), _dom_cla...
```

6.11 Communicate

```
[94]: # Get feature names from training data
features = X_train_over.columns
# Extract importances from model
importances = model.best_estimator_.named_steps['gradientboostingclassifier'].
    ↪feature_importances_
# Create a series with feature names and importances
feat_imp = pd.Series(importances,index=features).sort_values()
# Plot 10 most important features
feat_imp.tail(10).plot(kind='barh')
plt.xlabel("Gini Importance")
plt.ylabel("Feature")
plt.title("Feature Importance");
```



```
[95]: # save the model
import pickle
with open( "Gradient_boosting_model.pkl",'wb') as f :
    pickle.dump(model,f)
```

```
[96]: def make_predictions(data_filepath, model_filepath):
    # Wrangle JSON file
    X_test = pd.read_csv(data_filepath)
    # Load model
    with open(model_filepath,'rb') as file:
        model=pickle.load(file)
    # Generate predictions
    y_test_pred = model.predict(X_test)
```

```
# Put predictions into Series with name "bankrupt", and same index as X_test
y_test_pred = pd.Series(y_test_pred, index=X_test.index, name='bankrupt')
return y_test_pred
```

```
[97]: y_test_pred = make_predictions(
        data_filepath="data/X_test.csv",
        model_filepath="Gradient_boosting_model.pkl",
    )

    print("predictions shape:", y_test_pred.shape)
    y_test_pred.head()
```

predictions shape: (1364,)

```
[97]: 0    0
      1    0
      2    0
      3    1
      4    0
      Name: bankrupt, dtype: int64
```

```
[ ]:
```